MLOps in the Cloud: AWS vs. Azure vs. GCP Comparison (2026)

Behind every AI feature in a product lies a web of pipelines, registries, monitoring systems, and deployment infrastructure. That operational layer is what MLOps is about, and the three hyperscalers AWS, Azure and GCP have each built substantial platforms to address it.

This post breaks down the MLOps tooling offered by **Amazon Web Services**, **Microsoft Azure (Azure Machine Learning)**, and **Google Cloud Platform** covering core features, pricing models, and the honest pros and cons of each.

What is MLOps, and why does it matter?

MLOps (Machine Learning Operations) applies DevOps principles, version control, CI/CD, monitoring, automation to the ML lifecycle. It bridges the gap between data science experimentation and production-grade deployment, ensuring models are reproducible, auditable, scalable, and maintainable over time.

A mature MLOps platform typically covers:

- **Data management** ingestion, labelling, feature stores
- **Experiment tracking** logging runs, metrics, and hyperparameters
- **Pipeline orchestration** automating training, evaluation, and deployment
- **Model registry** versioning and staging models
- **CI/CD for ML** automating model promotion workflows
- **Monitoring** detecting drift, degradation, and data quality issues

All three cloud platforms address these areas, but they approach them very differently.

1. AWS Amazon

Launched in 2017, SageMaker is the most widely used managed ML platform on the market. It has grown into a sprawling suite of services under the “SageMaker Unified Studio” umbrella, attempting to unify data engineering, analytics, and ML into a single control plane.

Core MLOps Features

SageMaker Pipelines is the backbone of MLOps on AWS. It provides a directed acyclic graph (DAG) editor for defining end-to-end ML workflows data processing, training, evaluation, and deployment, with deep integration into EventBridge, Lambda, and CodePipeline for triggering runs based on S3 uploads or schedules.

SageMaker Projects adds CI/CD by providing infrastructure-as-code templates that provision standardized data scientist environments with source control, boilerplate code, and automated deployment pipelines. A/B testing and environment parity between dev and prod are built in.

SageMaker Model Monitor watches deployed endpoints for data drift, concept drift, bias, and quality degradation in real time. It integrates with SageMaker Clarify for explainability analysis.

SageMaker Model Registry manages model versions, metadata, approval states, and deployment lineage. Every step of a workflow is logged for full audit-trail reproducibility.

SageMaker Feature Store provides a managed repository for creating, sharing, and serving ML features, both online (low-latency) and offline (for batch training). It helps eliminate feature engineering duplication across teams.

SageMaker JumpStart gives access to a curated library of pre-trained foundation models, algorithms, and solution templates that can be deployed with a few clicks.

Pricing Model

SageMaker follows a **pay-as-you-go** model with no upfront fees or minimum commitments. However, the pricing is notoriously multi-dimensional:

- **Notebooks**: ~$0.05–$2/hour depending on instance type
- **Training compute**: GPU instances range from ~$0.50 to $30+/hour (on-demand); spot instances can reduce this by 60–90%
- **Inference endpoints**: ~$0.05–$5/hour depending on instance
- **Platform surcharge**: SageMaker's managed layer adds approximately 25–30% on top of raw EC2 instance costs (e.g., an `ml.p4d.24xlarge` at $42.07/hr vs. the underlying EC2 price of $32.77/hr)
- **MLOps-specific services** (Pipelines, Model Monitor, Feature Store, Data Wrangler) each carry separate charges

Savings options include SageMaker Savings Plans (1- or 3-year commitments) and Managed Spot Training, which makes GPU spot instances significantly more practical in training workflows.

Typical monthly spend for a team running continuous training and inference pipelines: **$1,000–$10,000+**, depending on instance class and usage intensity.

Pros

- Most mature and feature-rich managed ML platform available
- Deep integration with the entire AWS ecosystem (S3, IAM, VPC, EventBridge, Lambda, Redshift)
- Excellent enterprise security and compliance posture strong IAM, VPC isolation, audit logging
- SageMaker Managed Spot Training significantly reduces GPU training costs
- Broad framework support: TensorFlow, PyTorch, Hugging Face, MXNet, Scikit-learn, and more
- Kubeflow and custom container support for teams needing escape hatches

Cons

- Pricing is notoriously difficult to estimate in advance; month-end bill shock is common
- The platform surcharge (25–30% over raw EC2) is a real cost for teams with strong platform engineering capable of managing their own Kubernetes
- Steep learning curve — SageMaker's breadth means navigating dozens of overlapping services and APIs
- "Walled garden" architecture penalises multi-cloud strategies
- Debugging workflows and endpoints can be frustrating; tooling for local development and testing is limited
- UI, while improving, is not consistently intuitive

2. Microsoft Azure, Azure Machine Learning

Azure Machine Learning (Azure ML) is Microsoft's end-to-end platform for building, training, and deploying ML models. It is tightly woven into the Microsoft ecosystem Azure DevOps, Synapse Analytics, Power BI, Active Directory making it the natural choice for organisations already operating within the Microsoft stack.

Core MLOps Features

Azure ML Pipelines let teams compose reusable steps for data preparation, training, evaluation, and deployment. Pipelines integrate natively with Azure DevOps, allowing ML workflows to sit alongside software CI/CD in familiar tooling.

Model Registry in Azure ML provides versioning, staging, and lifecycle management. Users praise this as one of the platform's strongest features, model metadata, evaluation metrics, and lineage are captured consistently.

Responsible AI Dashboard consolidates explainability, fairness analysis, error analysis, and counterfactuals in a single interface, a clear area of differentiation as regulatory requirements around AI transparency grow.

Azure ML Designer provides a drag-and-drop visual canvas for building pipelines without writing code, lowering the barrier to entry for less technical users.

Model Monitoring tracks data drift and model performance in production, with configurable alerts. Integration with Azure Monitor and Application Insights enables operationally sophisticated setups.

Azure Arc integration enables hybrid and multi-cloud ML workloads models can be trained or deployed on-premise or on other clouds while still being managed through the Azure ML control plane. This is a meaningful differentiator for regulated industries with data residency requirements.

Azure ML CLI v2 and SDK v2 give code-first teams robust programmatic control over the entire platform, supporting YAML-based job definitions that integrate cleanly with GitOps workflows.

Pricing Model

Azure ML itself carries **no additional platform surcharge** the service layer is provided free. However, teams pay for every underlying Azure resource consumed:

- **Compute**: Billed at Azure VM rates; GPU clusters (e.g., NC A100 v4) carry approximately a **25% managed cluster surcharge** over raw VM on-demand prices (e.g., $39.91/hr managed vs. $31.93/hr raw VM)
- **Storage**: Azure Blob Storage, Azure Container Registry, Azure Key Vault, and Application Insights are all separately billed
- **Reserved Instances**: Commit to 1- or 3-year terms for up to 63% savings on underlying VMs; however, the Azure ML managed layer is not covered by reserved pricing
- **Azure Savings Plans**: Flexible commitment-based discounts across select compute

The lack of a direct SaaS fee is transparent in one sense, but the distributed billing across many Azure services makes **total cost estimation difficult in practice** a consistent complaint from users.

Typical monthly spend for a mid-sized team: similar to SageMaker, ranging from **$1,000–$10,000+**, with costs driven almost entirely by compute and storage consumption.

Pros

- Deep integration with Microsoft ecosystem (Azure DevOps, Synapse, Power BI, Entra ID / Azure AD)
- Strong enterprise governance: role-based access control, compliance certifications, and Responsible AI tooling
- Hybrid and multi-cloud capabilities via Azure Arc give flexibility SageMaker lacks
- No platform-level surcharge on the ML service itself
- Strong model registry and versioning frequently cited as best-in-class
- Low-code Designer interface broadens access to non-technical users
- Excellent support for regulated industries (healthcare, finance, government)

Cons

- Configuration overhead is significant: workspaces, IAM roles, networking, and identity must be carefully aligned before teams can move quickly
- Steeper onboarding curve, especially for teams new to the Azure ecosystem
- UI is polished but can obscure important settings behind multiple screens
- SDK compatibility issues have frustrated teams some users report being forced onto lower Python versions due to SDK incompatibilities
- AutoML has workflow limitations and restrictions on highly customised ML scenarios
- Costs, while transparent at the unit level, are very hard to forecast in aggregate

3. Google Cloud Platform, Vertex AI

Vertex AI, launched in 2021, is Google's answer to the fragmented tooling that preceded it (AutoML, AI Platform, and a constellation of separate APIs). The platform's philosophy is **abstraction over control** it aims to save you time rather than give you maximum infrastructure flexibility. In 2026, it has matured significantly and is particularly strong for teams building generative AI applications on top of Gemini.

Core MLOps Features

Vertex AI Pipelines are built on the open-source **Kubeflow Pipelines SDK**, making them portable and interoperable with other Kubernetes-native platforms. The DAG visualisation in the Google Cloud Console is widely considered the best in class cleaner and easier to debug than the AWS or Azure equivalents.

Vertex AI Model Registry supports versioning, aliases, and lifecycle management. It works across both AutoML and custom-trained models.

Vertex AI Feature Store enables teams to define, serve, and reuse ML features across training and serving, with online and offline serving paths.

Vertex AI Experiments tracks training runs, hyperparameters, metrics, and TensorBoard visualisations.

Vertex AI Model Monitoring watches production endpoints for data skew and model drift, with configurable alerts.

Vertex AI Workbench provides managed Jupyter notebooks deeply integrated with BigQuery and GCS, with collaborative capabilities for data science teams.

Model Garden gives access to 150+ models Gemini (Google's flagship), Imagen, Chirp, Veo, as well as third-party models (including Anthropic's Claude) and open-source models (Llama, Gemma). This breadth of foundation model access is a genuine advantage over SageMaker's Bedrock, which has its own curated model set.

Vertex AI Agent Builder (2025+) is a growing part of the platform aimed at enterprise agentic workflows, enabling teams to build, deploy, and govern multi-step AI agents grounded in enterprise data.

Pricing Model

Vertex AI uses a **pay-as-you-go** model billed per resource consumed:

- **Custom Training**: Billed per machine type per hour, based on Compute Engine rates, plus accelerator costs
- **AutoML**: Node-hour pricing importantly, deployed AutoML endpoints charge **continuously** even at zero traffic (a classification endpoint left running for 30 days costs ~$991 at $1.375/hour × 720 hours)
- **Managed Endpoints**: Billed by uptime; there is **no scale-to-zero** for standard endpoints, meaning idle endpoints accumulate costs
- **MLOps services** (Pipelines, Feature Store, Experiments): Usage-based, with some preview services offered free until general availability

- **Committed Use Discounts (CUDs)**: Offer up to 55% savings on predictable training workloads; one of the most impactful levers available
- **Free tier**: New Google Cloud accounts receive $300 in credits (90-day validity); certain preview services (Vertex AI Pipelines, TensorBoard, Gemini API) are free during preview

The **no scale-to-zero** behaviour for endpoints is the most significant hidden cost trap in the platform easily overlooked during initial deployment and capable of generating large unexpected bills.

Typical monthly spend: Comparable to the other platforms for similar workloads, but **highly variable** depending on endpoint management hygiene and training frequency.

Pros

- Kubeflow-based pipelines favour open standards and portability
- Best-in-class pipeline DAG visualisation easier to debug than AWS or Azure
- Deep, frictionless integration with BigQuery ML models can run directly inside BigQuery using SQL (`CREATE MODEL`), enabling analytics and ML in one place
- Model Garden is the most diverse multi-model marketplace, including both proprietary and open-source options
- Google's infrastructure advantage for TPU access (especially relevant for large model training)
- Committed Use Discounts can be very effective for predictable workloads
- Growing strength in generative AI and agentic workflows via Gemini

Cons

- No scale-to-zero for inference endpoints a significant and recurring source of runaway costs
- Complex, fragmented pricing that is difficult to forecast; no unified billing dashboard for the platform
- Steep learning curve, especially for teams unfamiliar with Google Cloud
- Documentation, while extensive, is frequently described as fragmented and hard to navigate
- Committed to the GCP ecosystem multi-cloud setups face friction
- Less mature enterprise governance tooling compared to Azure (though improving)
- AutoML endpoints charge continuously even when idle

Side-by-Side Comparison

Feature	AWS SageMaker	Azure ML	GCP Vertex AI
Pipeline orchestration	SageMaker Pipelines	Azure ML Pipelines + Azure DevOps	Vertex AI Pipelines (Kubeflow-based)
Model registry	SageMaker Model Registry	Azure ML Model Registry	Vertex AI Model Registry

Feature	AWS SageMaker	Azure ML	GCP Vertex AI
Feature store	SageMaker Feature Store	(Limited native support)	Vertex AI Feature Store
Monitoring	SageMaker Model Monitor + Clarify	Azure Monitor + ML Monitor	Vertex AI Model Monitoring
CI/CD integration	CodePipeline, GitHub Actions	Azure DevOps (native)	Cloud Build, GitHub Actions
Responsible AI tooling	SageMaker Clarify	Responsible AI Dashboard	Explainable AI
Foundation model access	Amazon Bedrock	Azure OpenAI Service	Model Garden (Gemini, Claude, Llama+)
Hybrid/multi-cloud	Limited	Azure Arc (strong)	Moderate
Platform pricing model	Pay-as-you-go + ~28% surcharge	No ML surcharge (compute billed separately)	Pay-as-you-go, node-hour billing
Scale-to-zero endpoints	Yes (serverless inference)	Yes	No
Open-source alignment	Moderate	Moderate	High (Kubeflow)
Best data integration	S3, Glue, Redshift	Synapse, Blob, Power BI	BigQuery, Dataflow, Looker
Typical monthly cost	$1,000–$10,000+	$1,000–$10,000+	$1,000–$10,000+

Pricing Summary

All three platforms use fundamentally similar approaches pay-as-you-go compute, with optional commitment-based discounts, but differ in important structural ways:

AWS SageMaker is the most expensive on a like-for-like compute basis due to its managed platform surcharge (approximately 25–30% over EC2 rates). Savings Plans and Managed Spot Training help significantly, but pricing complexity is the highest of the three.

Azure ML nominally has no platform-level fee, but the distributed billing across many Azure services makes total cost prediction difficult. The 25% managed cluster surcharge on GPU compute clusters is comparable to SageMaker's. Reserved Instances (1- or 3-year) on the underlying VMs offer up to 63% savings.

GCP Vertex AI has the most dangerous pricing trap idle endpoints that charge by the node-hour with no scale-to-zero, but Committed Use Discounts of up to 55% make it very competitive for teams with predictable, sustained workloads. The Kubeflow-based pipeline infrastructure also provides potential cost escape hatches by running training directly on GKE.

Which Platform Should You Choose?

The honest answer is that **your existing cloud footprint is the single biggest deciding factor**. All three platforms deliver comparable capability at comparable price points the differences are at the margins.

Choose AWS SageMaker if you are an AWS-native organisation with data in S3 and workloads on EC2/Lambda, you need maximum breadth of ML tooling, or you have a dedicated DevOps team capable of managing billing complexity. SageMaker is the default for MLOps roles and the most widely adopted platform in the industry.

Choose Azure ML if your organisation is Microsoft-centric (Microsoft 365, Azure DevOps, Azure AD), you operate in a regulated industry that requires strong governance and compliance tooling, or you need genuine hybrid/multi-cloud deployment via Azure Arc.

Choose GCP Vertex AI if your data already lives in BigQuery, you prefer open standards (Kubeflow), you are building primarily generative AI applications around Gemini, or your team values cleaner pipeline visualisation and a more developer-friendly abstraction. Budget carefully for endpoint idle time.

Final Thoughts

MLOps at scale is expensive regardless of the platform you choose. The real cost differential between providers is not in headline instance pricing, it is in engineering time spent fighting complexity, unexpected idle costs, and the labour of maintaining integrations. The platform that minimises friction for your specific team's stack will almost always be the most cost-effective in practice.

Before committing to any platform, run a realistic workload on the free tier or trial credits available from all three providers. The $300 in GCP credits, AWS Free Tier, and Azure's free trial are each substantial enough to prototype a real pipeline end-to-end.

Made by: Mate Kovasznai

Sources:

Amazon Web Services. (n.d.). *Amazon SageMaker documentation*. https://docs.aws.amazon.com/sagemaker/

Azure. (n.d.). *Azure Machine Learning documentation*. https://learn.microsoft.com/azure/machine-learning/

G2. (n.d.). *Amazon SageMaker reviews*. https://www.g2.com/

Gartner. (n.d.). *Gartner Peer Insights*. https://www.gartner.com/reviews

Google Cloud. (n.d.). *Vertex AI documentation*. https://cloud.google.com/vertex-ai/docs

nOps. (2026). *Vertex AI pricing guide*. https://www.nops.io/

Spendark. (n.d.). *Machine learning cost analysis*. https://spendark.com/

TrueFoundry. (n.d.). *Amazon SageMaker review*. https://www.truefoundry.com/